

Extra Credit Partial Sum

George Onwubuya

October 6, 2018

1 Main

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <math.h>
4  #include "Partial_Sum_Kernel.cu"
5
6  #define BLOCK_SIZE 32;
7  #define SAMPLE_SIZE 32
8
9  void FATAL (const char * s )
10 {
11     puts(s);
12     exit(1);
13 }
14
15 int main(int argc, char**argv) {
16
17     unsigned int data_size;
18     cudaError_t cuda_ret;
19
20
21     if(argc == 1) {
22         data_size= 64;
23     } else if(argc == 2) {
24         data_size= atoi(argv[1]);
25     } else {
26         printf("\n    Invalid input parameters!"
27             "\n    Usage: ./vecadd                # Vector of size
28             ↪ 10,000 is used"
29             "\n    Usage: ./vecadd <m>                # Vector of size
30             ↪ m is used"
31             "\n");
32         exit(0);
33     }
34 }
```

```

32
33 //Device data
34 int * array_dev;
35 int array_size = data_size;
36
37 //Host data
38 int * array_host = (int *) malloc (sizeof(int)*array_size);
39 for(int i = 0; i < data_size; ++i)
40     array_host[i] = i + 1;
41
42 for (int i = data_size; i < array_size; ++i)
43     array_host[i] = 0;
44
45 int expected_sum = data_size * (array_host[0] +
    ↪ array_host[data_size - 1]) / 2;
46
47 //Allocating & copying device memory
48 cuda_ret = cudaMalloc((void**)&array_dev,
    ↪ array_size*sizeof(int));
49     if(cuda_ret != cudaSuccess) FATAL("Unable to allocate
    ↪ device memory");
50 cuda_ret = cudaMemcpy(array_dev, array_host,
    ↪ array_size*sizeof(int), cudaMemcpyHostToDevice);
51     if(cuda_ret != cudaSuccess) FATAL("Unable to copy memory
    ↪ to device");
52
53 cudaDeviceSynchronize();
54
55 //Invoke Kernel}
56 vecSum_final_int1<<<dim3(1, 1, 1), dim3(SAMPLE_SIZE, 1,
    ↪ 1)>>>(array_dev);
57 vecSum_final_int<<<dim3(1, 1, 1), dim3(SAMPLE_SIZE, 1,
    ↪ 1)>>>(array_dev);
58
59 //Copying to host memory
60 int *result = (int *) malloc(sizeof(int)*array_size);
61 cuda_ret = cudaMemcpy(result, array_dev, sizeof(int)*array_size,
    ↪ cudaMemcpyDeviceToHost);
62 if(cuda_ret != cudaSuccess) FATAL("Unable to copy memory to
    ↪ host");
63
64 cudaDeviceSynchronize();
65
66 printf("Array size = %d\n", array_size);
67 printf("Expected result = %d\n", expected_sum);
68 printf("Calculated result = %d\n", result[0]);

```

```

69
70 for (int i = 0; i < data_size; ++i){
71     printf("[%2d] : %5d, %5d\n", i, array_host[i],
72         ↪ result[i]);
73 }
74
75 fflush(stdout);
76
77 free(array_host);
78 cudaFree(array_dev);
79
80 return 0;
81
82 };

```

2 Kernel

```

1  __global__ void vecSum_final_int(int * array)
2  {
3
4
5      for(unsigned int offset = blockDim.x; offset > 0; offset =
6         ↪ offset >> 1){
7          __syncthreads();
8
9          if (threadIdx.x < offset)
10             array[threadIdx.x] += array[threadIdx.x + offset];
11     }
12
13 __global__ void vecSum_final_int1(int * array)
14 {
15     const int tid = threadIdx.x << 1;
16
17     for (unsigned int stride = 1; stride <= blockDim.x; stride =
18         ↪ stride << 1 ){
19
20         __syncthreads();
21
22         if(tid % stride == 0)
23             array[tid] += array[tid + stride];
24     }
25 }

```

3 Output

```
==28325==NVPROF is profiling process 28325,command:/home/onwubuyag/Partial_Sum/partial_sum
==28325==Profiling application:/home/onwubuyag/Partial_Sum/partial_sum
==28325==Profiling result:
```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU:	50.92%	8.8960us	1	8.8960us	8.8960us	8.8960us	vecSum_final_int1(int*)
	27.29%	4.7680us	1	4.7680us	4.7680us	4.7680us	vecSum_final_int(int*)
	13.19%	2.3040us	1	2.3040us	2.3040us	2.3040us	[CUDA memcpy DtoH]
	8.60%	1.5030us	1	1.5030us	1.5030us	1.5030us	[CUDA memcpy HtoD]
API:	99.46%	145.53ms	1	145.53ms	145.53ms	145.53ms	cudaMalloc
	0.15%	219.99us	2	109.99us	15.702us	204.29us	cudaLaunch
	0.12%	181.70us	94	1.9330us	452ns	47.434us	cuDeviceGetAttribute
	0.11%	167.89us	1	167.89us	167.89us	167.89us	cuDeviceTotalMem
	0.08%	112.70us	1	112.70us	112.70us	112.70us	cudaFree
	0.04%	64.552us	2	32.276us	28.324us	36.228us	cudaMemcpy
	0.01%	15.640us	1	15.640us	15.640us	15.640us	cuDeviceGetName
	0.01%	15.548us	2	7.7740us	6.0780us	9.4700us	cudaDeviceSynchronize
	0.00%	6.1690us	3	2.0560us	565ns	4.2580us	cuDeviceGetCount
	0.00%	6.1590us	2	3.0790us	765ns	5.3940us	cudaSetupArgument
	0.00%	3.3440us	2	1.6720us	1.0260us	2.3180us	cudaConfigureCall
	0.00%	1.9730us	2	986ns	627ns	1.3460us	cuDeviceGet

4 Output Analysis

The output shows that the "vecSum final int" function has a faster execution time than the "vecSum final int1" function. This is down to the fact that the latter function reduces coalescing. The "vecSum final int1" has an execution time of 8.8960us and "vecSum final int" has an execution time of 4.7680us.