

Lab 1

George Onwubuya

September 5, 2018

1 Main File

```

/*****
*cr
*cr          (C) Copyright 2010 The Board of Trustees of the
*cr                      University of Illinois
*cr                      All Rights Reserved
*cr
*****/

#include <stdio.h>
#include "support.h"
#include "kernel.cu"

int main(int argc, char**argv) {

    Timer timer;
    cudaError_t cuda_ret;

    // Initialize host variables -----

    printf("\nSetting up the problem..."); fflush(stdout);
    startTime(&timer);

    unsigned int n;
    if(argc == 1) {
        n = 10000;
    } else if(argc == 2) {
        n = atoi(argv[1]);
    } else {
        printf("\n    Invalid input parameters!"
            "\n    Usage: ./vecadd          # Vector of size 10,000 is used"
            "\n    Usage: ./vecadd <m>      # Vector of size m is used"
            "\n");
        exit(0);
    }
}
```

```

}

float* A_h = (float*) malloc( sizeof(float)*n );
for (unsigned int i=0; i < n; i++) { A_h[i] = (rand()%100)/100.00 + 1; }

float* B_h = (float*) malloc( sizeof(float)*n );
for (unsigned int i=0; i < n; i++) { B_h[i] = (rand()%100)/100.00 + 1; }

float* C_h = (float*) malloc( sizeof(float)*n );

stopTime(&timer); printf("%f s\n", elapsedTime(timer));
printf("    Vector size = %u\n", n);

// Allocate device variables -----

printf("Allocating device variables..."); fflush(stdout);
startTime(&timer);

//INSERT CODE HERE
float* A_d;
cuda_ret = cudaMalloc((void**)&A_d, sizeof(float)*n );
    if(cuda_ret != cudaSuccess) FATAL("Unable to allocate device memory");

float* B_d;
cuda_ret = cudaMalloc((void**)&B_d, sizeof(float)*n );
    if(cuda_ret != cudaSuccess) FATAL("Unable to allocate device memory");

float* C_d;
cuda_ret = cudaMalloc((void**)&C_d, sizeof(float)*n );
    if(cuda_ret != cudaSuccess) FATAL("Unable to allocate device memory");

cudaDeviceSynchronize();
stopTime(&timer); printf("%f s\n", elapsedTime(timer));

// Copy host variables to device -----

printf("Copying data from host to device..."); fflush(stdout);
startTime(&timer);

//INSERT CODE HERE
cuda_ret = cudaMemcpy(A_d, A_h, sizeof(float)*n, cudaMemcpyHostToDevice);
    if(cuda_ret != cudaSuccess) FATAL("Unable to copy memory to device");

cuda_ret = cudaMemcpy(B_d, B_h, sizeof(float)*n, cudaMemcpyHostToDevice);
    if(cuda_ret != cudaSuccess) FATAL("Unable to copy memory to device");

```

```

cudaDeviceSynchronize();
stopTime(&timer); printf("%f s\n", elapsedTime(timer));

// Launch kernel -----

printf("Launching kernel..."); fflush(stdout);
startTime(&timer);

//INSERT CODE HERE
const unsigned int THREADS_PER_BLOCK = 256;
const unsigned int numBlocks = (n-1)/THREADS_PER_BLOCK + 1;
dim3 gridDim(numBlocks, 1, 1), blockDim(THREADS_PER_BLOCK, 1, 1);

vecAddKernel<<<numBlocks, THREADS_PER_BLOCK>>>(A_d, B_d, C_d, n);

cuda_ret = cudaDeviceSynchronize();
    if(cuda_ret != cudaSuccess) FATAL("Unable to launch kernel");
stopTime(&timer); printf("%f s\n", elapsedTime(timer));

// Copy device variables from host -----

printf("Copying data from device to host..."); fflush(stdout);
startTime(&timer);

//INSERT CODE HERE
cuda_ret = cudaMemcpy(C_h, C_d, sizeof(float)*n, cudaMemcpyDeviceToHost);
    if(cuda_ret != cudaSuccess) FATAL("Unable to copy memory from device");

cudaDeviceSynchronize();
stopTime(&timer); printf("%f s\n", elapsedTime(timer));

// Verify correctness -----

printf("Verifying results..."); fflush(stdout);

verify(A_h, B_h, C_h, n);

// Free memory -----

free(A_h);
free(B_h);
free(C_h);

//INSERT CODE HERE
cudaFree(A_d);

```

```

        cudaFree(B_d);
        cudaFree(C_d);

    return 0;

}

```

2 Kernel File

```

/*****
*cr
*cr          (C) Copyright 2010 The Board of Trustees of the
*cr                      University of Illinois
*cr                      All Rights Reserved
*cr
*****/

__global__ void vecAddKernel(float* A, float* B, float* C, int n) {

    // Calculate global thread index based on the block and thread indices ----

    //INSERT KERNEL CODE HERE
    int i = blockDim.x*blockIdx.x+threadIdx.x;

    // Use global index to determine which elements to read, add, and write ---

    //INSERT KERNEL CODE HERE

    if (i<n) C[i] = A[i] + B[i];

}

```

3 Output File

```

Setting up the problem...0.000044 s
    Vector size = 1000
Allocating device variables...0.398394 s
Copying data from host to device...0.012909 s
Launching kernel...0.006373 s
Copying data from device to host...0.000021 s
Verifying results...TEST PASSED

```

```

Setting up the problem...0.000246 s
    Vector size = 10000
Allocating device variables...0.375915 s
Copying data from host to device...0.006982 s
Launching kernel...0.002321 s
Copying data from device to host...0.000048 s
Verifying results...TEST PASSED

```

```

Setting up the problem...0.002300 s
    Vector size = 100000
Allocating device variables...0.377880 s
Copying data from host to device...0.000241 s
Launching kernel...0.002300 s
Copying data from device to host...0.000308 s
Verifying results...TEST PASSED

```

```

Setting up the problem...0.019750 s
    Vector size = 1000000
Allocating device variables...0.372675 s
Copying data from host to device...0.001425 s
Launching kernel...0.000158 s
Copying data from device to host...0.002088 s
Verifying results...TEST PASSED

```

4 Output Analysis

| Time Taken for Each Sub Process | | | |
|---------------------------------|---------------|---------------------|------------------|
| Elements(n) | Setting Up(s) | Device Variables(s) | Kernel Launch(s) |
| 1000 | 0.000044 | 0.398394 | 0.006373 |
| 10000 | 0.000246 | 0.375915 | 0.002321 |
| 100000 | 0.0023003 | 0.377880 | 0.002300 |
| 1000000 | 0.019750 | 0.372675 | 0.000158 |

| Time Taken for Each Sub Process | | |
|---------------------------------|-------------------|-------------------|
| Elements(n) | Host to Device(s) | Device to Host(s) |
| 1000 | 0.012909 | 0.000021 |
| 10000 | 0.006982 | 0.000048 |
| 100000 | 0.000241 | 0.000308 |
| 1000000 | 0.001425 | 0.002088 |