

Lab 3

George Onwubuya

September 27, 2018

1 Output Files

1.1 Output File 1

```
Setting up the problem...0.021003 s
  A: 1000 x 1000
  B: 1000 x 1000
  C: 1000 x 1000
Allocating device variables...0.173379 s
Copying data from host to device...0.004189 s
Launching kernel...0.008233 s
Copying data from device to host...0.004426 s
Verifying results...TEST PASSED
```

```
Setting up the problem...0.081349 s
  A: 2000 x 2000
  B: 2000 x 2000
  C: 2000 x 2000
Allocating device variables...0.182460 s
Copying data from host to device...0.017313 s
Launching kernel...0.060868 s
Copying data from device to host...0.012964 s
Verifying results...TEST PASSED
```

```
Setting up the problem...0.321950 s
  A: 4000 x 4000
  B: 4000 x 4000
  C: 4000 x 4000
Allocating device variables...0.186792 s
Copying data from host to device...0.068122 s
Launching kernel...0.485511 s
Copying data from device to host...0.045862 s
```

Verifying results...TEST PASSED

Setting up the problem...1.289725 s

A: 8000 x 8000

B: 8000 x 8000

C: 8000 x 8000

Allocating device variables...0.184046 s

Copying data from host to device...0.274219 s

Launching kernel...3.888804 s

Copying data from device to host...0.182651 s

Verifying results...TEST PASSED

1.2 Output File 2

Setting up the problem...0.009253 s

A: 1000 x 500

B: 500 x 500

C: 1000 x 500

Allocating device variables...0.205076 s

Copying data from host to device...0.001798 s

Launching kernel...0.002227 s

Copying data from device to host...0.002273 s

Verifying results...TEST PASSED

Setting up the problem...0.031236 s

A: 2000 x 1000

B: 1000 x 1000

C: 2000 x 1000

Allocating device variables...0.186218 s

Copying data from host to device...0.006258 s

Launching kernel...0.016270 s

Copying data from device to host...0.007320 s

Verifying results...TEST PASSED

Setting up the problem...0.119207 s

A: 4000 x 2000

B: 2000 x 2000

C: 4000 x 2000

Allocating device variables...0.187957 s

Copying data from host to device...0.024576 s

Launching kernel...0.121746 s

Copying data from device to host...0.022181 s
Verifying results...TEST PASSED

Setting up the problem...0.481023 s
A: 8000 x 4000
B: 4000 x 4000
C: 8000 x 4000
Allocating device variables...0.186625 s
Copying data from host to device...0.102642 s
Launching kernel...0.971854 s
Copying data from device to host...0.098304 s
Verifying results...TEST PASSED

2 Performance Analysis

2.1 Square Matrices (n x n)

Execution Time (seconds) for Each Process					
Elements(nxn)	Setting Up	DeviceVar	Kernel	HostToDevice	DeviceToHost
1000	0.021003	0.173379	0.008233	0.004189	0.004426
2000	0.081349	0.182460	0.060868	0.017313	0.012964
4000	0.321950	0.186792	0.485511	0.068122	0.045862
8000	1.289725	0.184046	3.888804	0.274219	0.182651

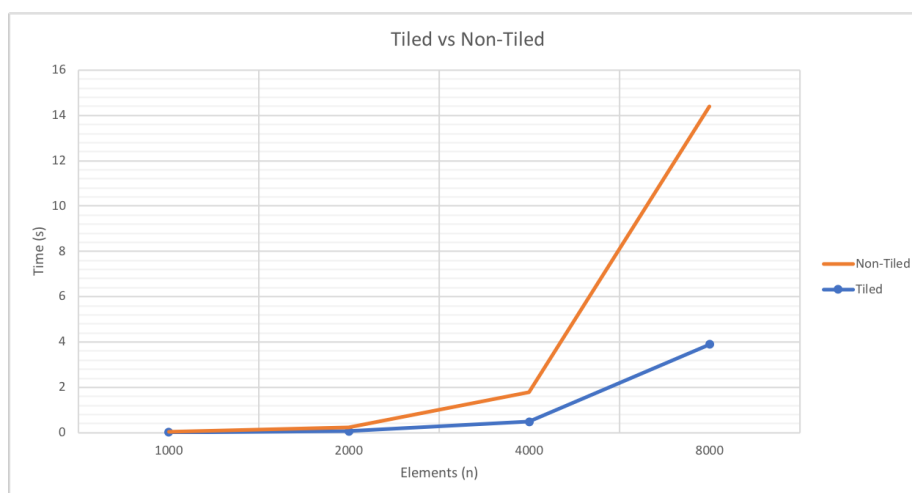
2.2 Rectangle Matrices (A = m x k and B = k x n)

Execution Time (seconds) for Each Process					
Elements (m,k,n)	Setting Up	DeviceVar	Kernel	HostToDevice	DeviceToHost
1000, 500, 500	0.009253	0.205076	0.002227	0.001798	0.002273
2000, 1000, 1000	0.031236	0.186218	0.016270	0.006258	0.007320
4000, 2000, 2000	0.119207	0.187957	0.121746	0.024576	0.022181
8000, 4000, 4000	0.481023	0.186625	0.971854	0.102642	0.098304

2.3 Square Matrix Comparison Between Tiled vs Non-Tiled

Kernel Execution Times (seconds)		
Elements (n x n)	Tiled	Non-Tiled
1000	0.008233	0.022178
2000	0.060868	0.163624
4000	0.485511	1.308468
8000	3.888804	10.493221

2.4 Tiled vs Non-Tiled Graphical Representation



2.5 Comments

For the square matrices each process time followed a similar pattern except for the process of allocating of 'device variables'. The time taken to allocate device variables is approximately the same regardless of the number of elements. As the number of elements increase the time taken setting up the problem, launching the kernel, copying data from the host to the device and vice versa also increases. There is an observable direct proportional relationship between these processes and the number of elements. The same conclusion can be made for the rectangular matrices.

There is an exponential increase in the kernel execution times associated with the tiled and non-tiled. However the non-tiled line graph shows a sharp increase in kernel execution time as the number of elements increase. This can be seen by the large spike between 4000 and 8000 elements on the graph. The tiled line graph does not exhibit that spike, as the purpose of tiling is reduce the number of global memory accesses. The graph shows that tiling reduces the kernel execution time by approximately a factor of 2.7.

3 Answers

3.1 C(i)

The execution of `nvcc -ptxas-options="-v" kernel.cu` revealed resource usage statistics which include 25 registers/thread, 2K (2048 bytes) of shared memory, 360 bytes of constant memory[0] and 4 bytes of constant memory[2]. There are obvious limits to the number of threads allocated to a single streaming processor as defined by hardware specifications. For the GeForce GTX 280 GPU that has a compute capability of 1.3 these specifications include: 512 threads/block, 16K shared memory/multiprocessor, 16K registers/multiprocessor, 8 blocks/multiprocessor, 1024 threads/multiprocessor and 124 registers/thread.

For this case, the block size is $16 \times 16 = 256$ threads and therefore the number of blocks/multiprocessor would be $1024 / 256 = 4$. However, $(256 \text{ threads} \times 25 \text{ reg/thread} \times 4 \text{ blocks/multiprocessor}) > 16\text{K registers/multiprocessor}$. The number of registers per multiprocessor is a limiting factor and therefore only 2 blocks can run on a single microprocessor which would yield $(256 \times 2 \times 25) < 16\text{K}$ and hence the total number of threads that can be simultaneously scheduled for 30 single multiprocessors would be $256 \times 2 \times 30 = 15360$.

4 Main

```
1  /*****
2  *cr
3  *cr          (C) Copyright 2010 The Board of Trustees of the
4  *cr          University of Illinois
5  *cr          All Rights Reserved
6  *cr
7
8  ↪  *****/
9  #include <stdio.h>
10 #include <stdlib.h>
11 #include "kernel.cu"
12 #include "support.h"
13
14 int main (int argc, char *argv[])
15 {
16
17     Timer timer;
18     cudaError_t cuda_ret;
```

```

19
20 // Initialize host variables
    ↪ -----
21
22 printf("\nSetting up the problem..."); fflush(stdout);
23 startTime(&timer);
24
25 float *A_h, *B_h, *C_h;
26 float *A_d, *B_d, *C_d;
27 size_t A_sz, B_sz, C_sz;
28 unsigned matArow, matAcol;
29 unsigned matBrow, matBcol;
30 dim3 dim_grid, dim_block;
31
32 if (argc == 1) {
33     matArow = 1000;
34     matAcol = matBrow = 1000;
35     matBcol = 1000;
36 } else if (argc == 2) {
37     matArow = atoi(argv[1]);
38     matAcol = matBrow = atoi(argv[1]);
39     matBcol = atoi(argv[1]);
40 } else if (argc == 4) {
41     matArow = atoi(argv[1]);
42     matAcol = matBrow = atoi(argv[2]);
43     matBcol = atoi(argv[3]);
44 } else {
45     printf("\n    Invalid input parameters!"
46           "\n    Usage: ./sgemm-tiled                # All matrices
    ↪    are 1000 x 1000"
47           "\n    Usage: ./sgemm-tiled <m>                # All matrices
    ↪    are m x m"
48           "\n    Usage: ./sgemm-tiled <m> <k> <n>        # A: m x k, B: k
    ↪    x n, C: m x n"
49           "\n");
50     exit(0);
51 }
52
53 A_sz = matArow*matAcol;
54 B_sz = matBrow*matBcol;
55 C_sz = matArow*matBcol;
56
57 A_h = (float*) malloc( sizeof(float)*A_sz );
58 for (unsigned int i=0; i < A_sz; i++) { A_h[i] =
    ↪ (rand()%100)/100.00; }
59

```

```

60     B_h = (float*) malloc( sizeof(float)*B_sz );
61     for (unsigned int i=0; i < B_sz; i++) { B_h[i] =
        ↪ (rand()%100)/100.00; }
62
63     C_h = (float*) malloc( sizeof(float)*C_sz );
64
65     stopTime(&timer); printf("%f s\n", elapsedTime(timer));
66     printf("      A: %u x %u\n      B: %u x %u\n      C: %u x %u\n",
        ↪ matArow, matAcol,
67           matBrow, matBcol, matArow, matBcol);
68
69     // Allocate device variables
        ↪ -----
70
71     printf("Allocating device variables..."); fflush(stdout);
72     startTime(&timer);
73
74     //INSERT CODE HERE
75     cuda_ret = cudaMalloc((void**)&A_d, sizeof(float)*A_sz);
76     if (cuda_ret != cudaSuccess) FATAL("Unable to
        ↪ allocate device memory");
77
78     cuda_ret = cudaMalloc((void**)&B_d, sizeof(float)*B_sz);
79     if (cuda_ret != cudaSuccess) FATAL("Unable to allocate
        ↪ device memory");
80
81     cuda_ret = cudaMalloc((void**)&C_d, sizeof(float)*C_sz);
82     if (cuda_ret != cudaSuccess) FATAL("Unable to allocate
        ↪ device memory");
83
84     cudaDeviceSynchronize();
85     stopTime(&timer); printf("%f s\n", elapsedTime(timer));
86
87     // Copy host variables to device
        ↪ -----
88
89     printf("Copying data from host to device...");
90     ↪ fflush(stdout);
91     startTime(&timer);
92
93     //INSERT CODE HERE
94     cuda_ret = cudaMemcpy(A_d, A_h, sizeof(float)*A_sz,
        ↪ cudaMemcpyHostToDevice);
95     if (cuda_ret != cudaSuccess) FATAL("Unable to copy to
        ↪ device memory");

```

```

96     cuda_ret = cudaMemcpy(B_d, B_h, sizeof(float)*B_sz,
    ↪     cudaMemcpyHostToDevice);
97     if (cuda_ret != cudaSuccess) FATAL("Unable to copy to
    ↪     device memoy");
98
99
100    cudaDeviceSynchronize();
101    stopTime(&timer); printf("%f s\n", elapsedTime(timer));
102
103    // Launch kernel using standard sgemm interface
    ↪    -----
104    printf("Launching kernel..."); fflush(stdout);
105    startTime(&timer);
106    basicSgemm('N', 'N', matArow, matBcol, matBrow, 1.0f, \
107              A_d, matArow, B_d, matBrow, 0.0f, C_d, matBrow);
108
109    cuda_ret = cudaDeviceSynchronize();
110    if(cuda_ret != cudaSuccess) FATAL("Unable to launch
    ↪    kernel");
111    stopTime(&timer); printf("%f s\n", elapsedTime(timer));
112
113    // Copy device variables from host
    ↪    -----
114
115    printf("Copying data from device to host...");
    ↪    fflush(stdout);
116    startTime(&timer);
117
118    //INSERT CODE HERE
119    cuda_ret = cudaMemcpy(C_h, C_d, sizeof(float)*C_sz,
    ↪    cudaMemcpyDeviceToHost);
120    if (cuda_ret != cudaSuccess) FATAL("Unable to copy to
    ↪    host memory");
121
122    cudaDeviceSynchronize();
123    stopTime(&timer); printf("%f s\n", elapsedTime(timer));
124
125    // Verify correctness
    ↪    -----
126
127    printf("Verifying results..."); fflush(stdout);
128
129    verify(A_h, B_h, C_h, matArow, matAcol, matBcol);
130
131

```



```

132 // Free memory
    ↪ -----
133
134 free(A_h);
135 free(B_h);
136 free(C_h);
137
138 //INSERT CODE HERE
139 cudaFree(A_d);
140 cudaFree(B_d);
141 cudaFree(C_d);
142 return 0;
143
144 }

```

5 Kernel

```

1  /*****
2  *cr
3  *cr          (C) Copyright 2010 The Board of Trustees of the
4  *cr          University of Illinois
5  *cr          All Rights Reserved
6  *cr
7
    ↪  *****/
8
9  #include <stdio.h>
10
11 #define TILE_SZ 16
12
13 __global__ void mysgemm(int m, int n, int k, const float *A,
    ↪  const float *B, float* C) {
14
15
16     ↪  /*****
17     *
18     * Compute C = A x B
19     * where A is a (m x k) matrix
20     * where B is a (k x n) matrix
21     * where C is a (m x n) matrix
22     *
23     * Use shared memory for tiling
24     *

```

```

24 ↪ *****/
25
26 // INSERT KERNEL CODE HERE
27 unsigned int TiRow = threadIdx.y;
28 unsigned int TiCol = threadIdx.x;
29 unsigned int row = blockIdx.y * blockDim.y + threadIdx.y;
30 unsigned int col = blockIdx.x * blockDim.x + threadIdx.x;
31
32 __shared__ float As[TILE_SZ][TILE_SZ];
33 __shared__ float Bs[TILE_SZ][TILE_SZ];
34
35 float sum = 0;
36
37
38 for(unsigned int TiNum = 0; TiNum < (k-1)/TILE_SZ+1;
39 ↪ TiNum++){
40     if((row < m) && (TiNum * TILE_SZ + TiCol) < k)
41         As[TiRow][TiCol] = A[row * k + TiNum * TILE_SZ
42 ↪ + TiCol];
43     else
44         As[TiRow][TiCol] = 0;
45
46     if((TiNum * TILE_SZ + TiRow) < k && col < n)
47         Bs[TiRow][TiCol] = B[(TiNum * TILE_SZ +
48 ↪ TiRow) * n + col];
49     else
50         Bs[TiRow][TiCol] = 0;
51     __syncthreads();
52
53     //Calculate inner product for the tile
54     //Checking for matrix size to lower power and practice
55     ↪ green computing
56     if(row < m && col < n)
57         for(unsigned int TiElem = 0; TiElem < TILE_SZ;
58 ↪ TiElem++)
59             sum = sum +
60             ↪ As[TiRow][TiElem]*Bs[TiElem][TiCol];
61     __syncthreads();
62 }
63
64 //Prevent writing of output to an undefined block
65 if (row < m && col < n)
66     C[row * n + col] = sum;
67 }

```

```

63
64 void basicSgemm(char transa, char transb, int m, int n, int k,
    ↪ float alpha, const float *A, int lda, const float *B, int
    ↪ ldb, float beta, float *C, int ldc)
65 {
66     if ((transa != 'N') && (transa != 'n')) {
67         printf("unsupported value of 'transa'\n");
68         return;
69     }
70
71     if ((transb != 'N') && (transb != 'n')) {
72         printf("unsupported value of 'transb'\n");
73         return;
74     }
75
76     if ((alpha - 1.0f > 1e-10) || (alpha - 1.0f < -1e-10)) {
77         printf("unsupported value of alpha\n");
78         return;
79     }
80
81     if ((beta - 0.0f > 1e-10) || (beta - 0.0f < -1e-10)) {
82         printf("unsupported value of beta\n");
83         return;
84     }
85
86     // Initialize thread block and kernel grid dimensions
    ↪ -----
87
88     //INSERT CODE HERE
89     const unsigned int BLOCK_SIZE = TILE_SZ; //use 16 x 16 thread
    ↪ blocks
90
91     dim3 block(BLOCK_SIZE, BLOCK_SIZE, 1);
92     dim3 grid((n + BLOCK_SIZE - 1)/BLOCK_SIZE, (m + BLOCK_SIZE
    ↪ -1)/BLOCK_SIZE, 1);
93
94
95     // Invoke CUDA kernel
    ↪ -----
96
97     //INSERT CODE HERE
98
99     mysgemm<<< grid, block>>>(m, n, k, A, B, C);
100
101 }

```