

# Lab 6

George Onwubuya

November 15, 2018

## 1 Histogram

### 1.1 Output

```
Setting up the problem...0.002114 s
  Input size = 100000
  Number of bins = 4096
Allocating device variables...0.000280 s
Copying data from host to device...0.000272 s
Launching kernel...0.000452 s
Copying data from device to host...0.000039 s
Verifying results...TEST PASSED
```

```
Setting up the problem...0.003954 s
  Input size = 200000
  Number of bins = 4096
Allocating device variables...0.000304 s
Copying data from host to device...0.000484 s
Launching kernel...0.000466 s
Copying data from device to host...0.000040 s
Verifying results...TEST PASSED
```

```
Setting up the problem...0.008680 s
  Input size = 400000
  Number of bins = 4096
Allocating device variables...0.000328 s
Copying data from host to device...0.000805 s
Launching kernel...0.000458 s
Copying data from device to host...0.000039 s
Verifying results...TEST PASSED
```

```

Setting up the problem...0.017020 s
    Input size = 800000
    Number of bins = 4096
Allocating device variables...0.000258 s
Copying data from host to device...0.001406 s
Launching kernel...0.000510 s
Copying data from device to host...0.000043 s
Verifying results...TEST PASSED

```

```

Setting up the problem...0.022502 s
    Input size = 1000000
    Number of bins = 4096
Allocating device variables...0.000264 s
Copying data from host to device...0.001899 s
Launching kernel...0.000525 s
Copying data from device to host...0.000041 s
Verifying results...TEST PASSED

```

```

Setting up the problem...0.033799 s
    Input size = 1600000
    Number of bins = 4096
Allocating device variables...0.000311 s
Copying data from host to device...0.002794 s
Launching kernel...0.000543 s
Copying data from device to host...0.000040 s
Verifying results...TEST PASSED

```

## 1.2 Performance Analysis (4096 Bins)

Execution Time (seconds) for Each Process					
Elements(m)	Setting Up	DeviceVar	HostToDevice	Kernel	DeviceToHost
100000	0.002114	0.000280	0.000272	0.000452	0.000039
200000	0.003954	0.000304	0.000484	0.000466	0.000040
400000	0.008680	0.000328	0.000805	0.000458	0.000039
800000	0.017020	0.000258	0.001406	0.000510	0.000043
1000000	0.022502	0.000264	0.001899	0.000525	0.000041
1600000	0.033799	0.000311	0.002794	0.000543	0.000040

### 1.2.1 Comments

When setting up the shell script file, the number of elements were varied but I chose to use the same number of bins, 4096. The time taken to allocate device variables and copy from device to host are approximately the same. This is

because the same number of variables are allocated regardless of the size of the array of input elements. The times taken to set up the problem, copy from host to device and launching the kernel are all directly proportional to the size of the input array.

## 2 Histogram (Optimized)

### 2.1 Output

```
Setting up the problem...0.002145 s
    Input size = 100000
    Number of bins = 4096
Allocating device variables...0.000262 s
Copying data from host to device...0.000236 s
Launching kernel...0.000498 s
Copying data from device to host...0.000036 s
Verifying results...TEST PASSED
```

```
Setting up the problem...0.004105 s
    Input size = 200000
    Number of bins = 4096
Allocating device variables...0.000303 s
Copying data from host to device...0.000387 s
Launching kernel...0.000504 s
Copying data from device to host...0.000042 s
Verifying results...TEST PASSED
```

```
Setting up the problem...0.008779 s
    Input size = 400000
    Number of bins = 4096
Allocating device variables...0.000259 s
Copying data from host to device...0.000713 s
Launching kernel...0.000528 s
Copying data from device to host...0.000041 s
Verifying results...TEST PASSED
```

```
Setting up the problem...0.016875 s
    Input size = 800000
    Number of bins = 4096
Allocating device variables...0.000261 s
Copying data from host to device...0.001438 s
```

```

Launching kernel...0.000560 s
Copying data from device to host...0.000041 s
Verifying results...TEST PASSED

```

```

Setting up the problem...0.021037 s
    Input size = 1000000
    Number of bins = 4096
Allocating device variables...0.000290 s
Copying data from host to device...0.001656 s
Launching kernel...0.000589 s
Copying data from device to host...0.000040 s
Verifying results...TEST PASSED

```

```

Setting up the problem...0.033860 s
    Input size = 1600000
    Number of bins = 4096
Allocating device variables...0.000284 s
Copying data from host to device...0.002730 s
Launching kernel...0.000629 s
Copying data from device to host...0.000040 s
Verifying results...TEST PASSED

```

## 2.2 Performance Analysis

### 2.2.1 Array Size

Execution Time (seconds) for Each Process					
Elements(m)	Setting Up	DeviceVar	HostToDevice	Kernel	DeviceToHost
100000	0.002145	0.000262	0.000236	0.000498	0.000036
200000	0.004105	0.000303	0.000387	0.000504	0.000042
400000	0.008779	0.000259	0.000713	0.000528	0.000041
800000	0.016875	0.000261	0.001438	0.000560	0.000041
1000000	0.021037	0.000290	0.001656	0.000589	0.000040
1600000	0.033860	0.000284	0.002730	0.000629	0.000040

### 2.2.2 Comments

The execution times relate to the same way to the initial histogram kernel that had no optimization.

## 2.3 Answers

### 2.3.1 Description of Optimization

To optimize the histogram kernel, the kernel code was written such that a location in shared memory would only be updated once. Three new variables were introduced to the kernel code. The first two variables current index and previous index were meant to keep track of the elements from the input array. The current index and previous were compared to check if the same element was still being evaluated because of possible repetitions in data. If they were found to be the same, a third variable accumulator would store the number of times this element occurred in succession in the array . However if the elements were not the same, the value in the accumulator would be written to a location in shared memory as indicated by the previous index, the previous index would then be set to the current index and the value in the accumulator would be set to one to indicate a shift to the next element. The accumulator and the previous index variables were initialized to zero outside the while loop and the current index variable was set to the current input element under evaluation by the histogram kernel code.

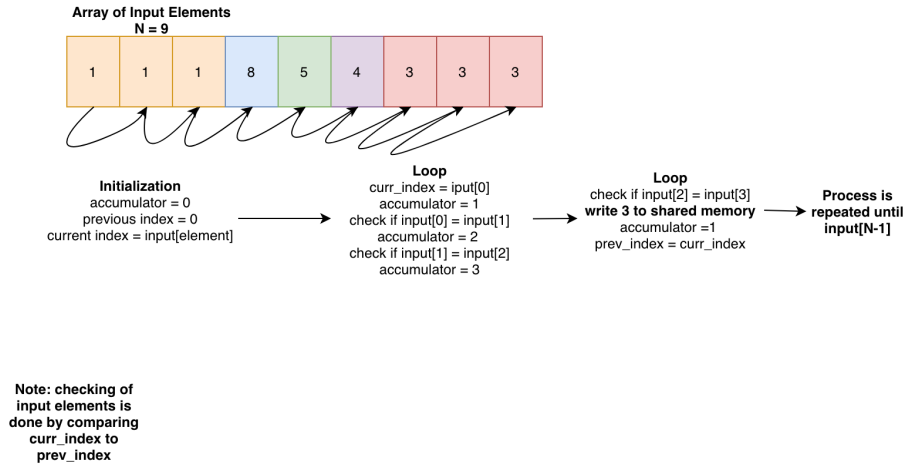


Figure 1: An illustration of what the kernel code does

### 2.3.2 Difficulties with Optimization

A couple of difficulties I experienced during this optimization were found in not analyzing the the second while loop structure of the code that is used to update the value in the accumulator to the shared memory location. I examined how the loop would behave at the beginning of the and during the middle part of the execution and neglected to check what happens when the loop ends. As a result my code was incomplete and was unable to simulate past the first bin. I also made the mistake of attempting to carry out the optimization without first checking if the code was properly written out.

### 2.3.3 Change in Execution Time

Kernel Times in microseconds		
Elements(m)	Non-Optimized	Optimized
100000	11.776	76.414
200000	18.367	82.398
400000	32.614	99.646
800000	63.039	129.12
1000000	77.356	144.09
1600000	120.38	190.17

```

==16836== NVPROF is profiling process 16836, command: /home/onwubuyag/mplabs/lab6-histogram/histogram 1600000
==16836== Profiling application: /home/onwubuyag/mplabs/lab6-histogram/histogram 1600000
==16836== Profiling result:

```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	94.55%	2.3343ms	1	2.3343ms	2.3343ms	2.3343ms	[CUDA memcpy HtoD]
	4.88%	120.38us	1	120.38us	120.38us	120.38us	histogram_kernel(unsigned int*, unsigned int*, unsigned int, unsigned int)
	0.28%	6.9760us	2	3.4880us	2.2720us	4.7040us	[CUDA memset]
	0.17%	4.3200us	1	4.3200us	4.3200us	4.3200us	[CUDA memcpy DtoH]
	0.11%	2.7520us	1	2.7520us	2.7520us	2.7520us	convert_kernel(unsigned int*, unsigned char*, unsigned int)
API calls:	97.41%	141.24ms	4	35.310ms	102.12us	140.26ms	cudaFree
	1.90%	2.7523ms	2	1.3762ms	33.967us	2.7183ms	cudaMemcpy
	0.27%	392.60us	3	130.87us	118.50us	155.02us	cudaMalloc
	0.11%	157.60us	2	78.799us	15.451us	142.15us	cudaLaunch
	0.11%	154.40us	94	1.6420us	451ns	38.678us	cuDeviceGetAttribute
	0.10%	149.68us	1	149.68us	149.68us	149.68us	cuDeviceTotalMem
	0.05%	70.629us	4	17.657us	4.6120us	47.116us	cudaDeviceSynchronize
	0.02%	28.445us	2	14.222us	11.024us	17.421us	cudaMemset
	0.02%	23.885us	1	23.885us	23.885us	23.885us	cuDeviceGetName
	0.01%	9.9800us	7	1.4250us	439ns	6.4630us	cudaSetupArgument
	0.00%	6.1310us	3	2.0430us	585ns	4.3260us	cuDeviceGetCount
	0.00%	2.6600us	2	1.3300us	887ns	1.7730us	cudaConfigureCall
	0.00%	1.7650us	2	882ns	608ns	1.1570us	cuDeviceGet

Figure 2: Nvidia profiler result for non-optimized histogram kernel for m = 1600000

```

==17003== NVPROF is profiling process 17003, command: /home/onwubuyag/mpplabs/lab6-histogram/histogram 1600000
==17003== Profiling application: /home/onwubuyag/mpplabs/lab6-histogram/histogram 1600000
==17003== Profiling result:

```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	91.79%	2.2817ms	1	2.2817ms	2.2817ms	2.2817ms	[CUDA memcpy HtoD]
	7.65%	190.17us	1	190.17us	190.17us	190.17us	histogram_kernel_optimized(unsigned int*, unsigned int*, unsigned int, unsigned int)
	0.28%	6.8800us	2	3.4400us	2.2400us	4.6400us	[CUDA memset]
	0.18%	4.4800us	1	4.4800us	4.4800us	4.4800us	[CUDA memcpy DtoH]
	0.11%	2.6560us	1	2.6560us	2.6560us	2.6560us	convert_kernel(unsigned int*, unsigned char*, unsigned int)
API calls:	97.40%	142.05ms	4	35.513ms	104.83us	140.98ms	cudaFree
	1.84%	2.6820ms	2	1.3410ms	33.133us	2.648ms	cudaMemcpy
	0.77%	387.24us	3	129.08us	103.39us	151.88us	cudaMalloc
	0.14%	205.15us	1	205.15us	205.15us	205.15us	cuDeviceTotalMem
	0.14%	202.01us	94	2.1490us	442ns	56.735us	cuDeviceGetAttribute
	0.11%	159.70us	2	79.850us	16.012us	143.69us	cudaLaunch
	0.05%	76.425us	4	19.106us	4.9700us	51.634us	cudaDeviceSynchronize
	0.02%	29.115us	2	14.557us	10.940us	18.166us	cudaMemset
	0.02%	28.463us	1	28.463us	28.463us	28.463us	cuDeviceGetName
	0.01%	10.258us	7	1.4650us	452ns	6.7510us	cudaSetupArgument
	0.00%	5.9090us	3	1.9690us	527ns	4.3270us	cuDeviceGetCount
	0.00%	2.7080us	2	1.3540us	907ns	1.8810us	cudaConfigureCall
	0.00%	1.6430us	2	821ns	545ns	1.0980us	cuDeviceGet

Figure 3: Nvidia profiler result for optimized histogram kernel for  $m = 1600000$

### 2.3.4 Explanation for Optimization

The optimization increased the amount of time spent executing the kernel. I suspect that might be the result of the input array being generated with less repetition than originally anticipated. The contention rate was low and hence the addition of more variables would only serve to increase the execution time.

## 2.4 Kernel

```

1  /*****
2  *cr
3  *cr          (C) Copyright 2010 The Board of Trustees of the
4  *cr          University of Illinois
5  *cr          All Rights Reserved
6  *cr
7
8  →  *****/
9
10 // Define your kernels in this file you may use more than one
11 → kernel if you
12 // need to
13
14 __global__ void histogram_kernel(unsigned int* input, unsigned
15 → int* bins,
16 → unsigned int num_elements, unsigned int num_bins){
17
18     extern __shared__ unsigned int bins_s[];
19
20     //Shared Memory
21     int thid = threadIdx.x;
22     while(thid < num_bins){
23
24         bins_s[thid] = 0u;
25         thid += blockDim.x;

```

```

22     }
23     __syncthreads();
24
25
26     //Histogram calculation
27     unsigned int element = blockIdx.x * blockDim.x +
        ↪ threadIdx.x;
28
29     while(element < num_elements){
30
31         atomicAdd(&(bins_s[input[element]]), 1);
32         element += blockDim.x * gridDim.x;
33     }
34     __syncthreads();
35
36     //Global Memory
37     thid = threadIdx.x;
38     while(thid < num_bins){
39
40         atomicAdd(&(bins[thid]), bins_s[thid]);
41         thid += blockDim.x;
42     }
43 }
44
45
46 __global__ void histogram_kernel_optimized(unsigned int* input,
    ↪ unsigned int* bins,
47     unsigned int num_elements, unsigned int num_bins) {
48
49     // INSERT CODE HERE
50     extern __shared__ unsigned int bins_s[];
51
52     //Shared memory
53     int thid = threadIdx.x;
54     while ( thid < num_bins){
55
56         bins_s[thid] = 0u;
57         thid += blockDim.x;
58     }
59     __syncthreads();
60
61     //Histogram calculation
62     unsigned int element = blockIdx.x * blockDim.x +
        ↪ threadIdx.x;
63     unsigned int accumulator = 0;
64     unsigned int prev_index = 0;

```



```

65
66     while(element < num_elements){
67
68         unsigned int curr_index = input[element];
69
70         if(curr_index != prev_index){
71
72             atomicAdd(&(bins_s[prev_index]),
73                 ↪ accumulator);
74             accumulator = 1;
75             prev_index = curr_index;
76
77         }
78
79         else{
80             accumulator++;
81
82         }
83         element += blockDim.x * gridDim.x;
84     }
85     if(accumulator > 0){
86         atomicAdd(&(bins_s[prev_index]), accumulator);
87     }
88     __syncthreads();
89
90     //Global memory
91     thid = threadIdx.x;
92     while(thid < num_bins){
93
94         atomicAdd(&(bins[thid]), bins_s[thid]);
95         thid += blockDim.x;
96     }
97 }
98
99 __global__ void convert_kernel(unsigned int *bins32, uint8_t
100 ↪ *bins8,
101     unsigned int num_bins) {
102
103     // INSERT CODE HERE
104     int thid = blockIdx.x * blockDim.x + threadIdx.x;
105
106     while (thid < num_bins){
107
108         //Use local register value (avoids copying from
109         ↪ global twice)
110         unsigned int reg_bin = bins32[thid];

```

```

108
109         if(reg_bin > 255){
110             bins8[thid] = 255u;
111         }
112
113         else{
114             bins8[thid] = (uint8_t) reg_bin;
115         }
116         thid += blockDim.x * gridDim.x;
117     }
118
119 }
120
121 *****
122 Setup and invoke your kernel(s) in this function. You may also
123 ↪ allocate more
124 GPU memory if you need to
125 *****
126 void histogram(unsigned int* input, uint8_t* bins, unsigned int
127 ↪ num_elements,
128             unsigned int num_bins) {
129
130     // Create 32 bit bins
131     unsigned int *bins32;
132     cudaMalloc((void*)&bins32, num_bins * sizeof(unsigned int));
133     cudaMemset(bins32, 0, num_bins * sizeof(unsigned int));
134
135     // Launch histogram kernel using 32-bit bins
136     dim3 dim_grid, dim_block;
137     dim_block.x = 512; dim_block.y = dim_block.z = 1;
138     dim_grid.x = 30; dim_grid.y = dim_grid.z = 1;
139
140     //Comment out the kernel not used
141     //histogram_kernel<<<dim_grid, dim_block,
142     ↪ num_bins*sizeof(unsigned int)>>>
143     // (input, bins32, num_elements, num_bins);
144     histogram_kernel_optimized<<<dim_grid, dim_block,
145     ↪ num_bins*sizeof(unsigned int)>>>
146     (input, bins32, num_elements, num_bins);
147
148     // Convert 32-bit bins into 8-bit bins
149     dim_block.x = 512;
150     dim_grid.x = (num_bins - 1)/dim_block.x + 1;
151     convert_kernel<<<dim_grid, dim_block>>>(bins32, bins,
152     ↪ num_bins);
153
154 }

```

```
149      // Free allocated device memory  
150      cudaFree(bins32);  
151  
152  }
```